



Thes1s

Multi-Agent Orchestration Alternatives & Model Optimization

Evaluating 11 platforms for the Pitch Deck pipeline blocked on Anthropic Managed Agents

April 20, 2026

Thes1s Project — Internal Research Report

Decision-support for the Pitch Deck pipeline



Table of Contents

1. Executive Summary
2. The Problem We Are Solving
3. Current State of the Pipeline
4. Evaluation Framework
5. Platform Deep Dives
6. Comparison Matrix
7. Migration Effort by Platform
8. Cost Analysis
9. Model Optimization Per Agent
10. Production Patterns We Are Missing
11. Recommended Path Forward
12. Appendix — Agent Inventory



1. Executive Summary

You are 1.5 weeks into waiting for Anthropic's multiagent Research Preview approval. The One Pager pipeline is live. The Pitch Deck (10 specialists + coordinator across 5 dependency waves) and Full Story (7 agents across 2 phases) cannot ship without orchestrator-to-subagent dispatch. This report evaluates 11 alternative platforms across four categories, scores them against your specific constraints, and recommends a phased path forward.

Top-Line Recommendations

1. Primary path (~5 day migration): Move the Pitch Deck coordinator to Claude Agent SDK (Anthropic's open-source self-hosted agent loop). Subagent files are nearly identical to your current `prompt.md` + `managed-agent.yaml` format. Production parity preserved — when `Managed Agents callable_agents` opens, you can swap back without prompt changes.
2. Backup path if Claude Agent SDK is too immature: Inngest AgentKit. AI-first multi-agent framework on a durable execution backend. Free tier covers your volume. Wave structure maps to Network + Router pattern almost directly.
3. Long-term Cloudflare-native answer: Cloudflare Project Think (sub-agents on Durable Objects). Currently in preview. Once GA, this is the right home for Thes1s — same vendor as your Workers + D1 + R2 stack.
4. Model optimization is independent of platform choice and worth doing now. Mixing models per agent — Gemini 2.5 Pro for filing reading, GPT-5.4 for valuation math, Claude Opus 4.7 for synthesis and bear-case reasoning, Claude Haiku 4.5 for orchestration — saves ~40% per Pitch Deck and lifts quality on numerical sections.
5. What we are missing today: real observability (LangSmith / Langfuse / AgentOps trace + replay), structured eval harness against known verdicts, prompt versioning with measured impact, circuit breakers on cost + token spend, and typed structured output validation at agent boundaries.



2. The Problem We Are Solving

Your `feedback_no_custom_orchestration` memory is the load-bearing constraint: Workers, Durable Objects, and hand-rolled dispatch all failed in v1 of the pipeline. The pivot to Managed Agents was the right call. But "no custom orchestration" was always shorthand for "no DIY orchestration code I have to debug myself" — it was never a vote against using a battle-tested orchestration framework someone else maintains.

The Two Functional Requirements

- Multi-agent callability with parallelism within waves. The Pitch Deck dispatches 10 specialist agents in 5 dependency waves. Within Wave 2, three agents (Moats, Financial Analyst, Management Evaluator) must run in parallel. Wave 3 cannot start until Wave 2 completes.
- A smart, dynamic orchestrator that handles failures. Pitch Deck and Full Story runs take 10-30 minutes. Network blips, rate-limit 429s, malformed JSON, and partial section outputs happen in production. The orchestrator must retry transient failures, surface durable failures, and continue the pipeline when individual sections fail rather than aborting the whole 25-minute run.

Anything you adopt must satisfy both, and must not violate `feedback_production_parity` — the migration must be reversible if Managed Agents access lands tomorrow.




3. Current State of the Pipeline

Pitch Deck Wave Structure

The current pipeline is a 5-superstep DAG. Every modern multi-agent framework has a primitive for this — the question is which one we want to own.

Pitch Deck — 5 Wave Dependency Graph



Current Model Assignments (20 agents)

19 of 20 agents are on Sonnet 4.6. Only risk-analyst-fullstory runs Opus. Total prompt corpus is ~98,000 words across 20 agents — roughly 130K tokens of system prompt content. This is a meaningful asset; any migration must preserve it byte-for-byte.



Agent	Stage	Current Model	Prompt Words
coordinator-pitchdeck	Pitch Deck	Sonnet 4.6	1,110
coordinator-fullstory	Full Story	Sonnet 4.6	1,926
one-pager	One Pager	Sonnet 4.6	2,625
annual-reader	PSR (PD+FS)	Sonnet 4.6	5,424
quarterly-reader	PSR (PD+FS)	Sonnet 4.6	4,424
business-analyst-pitchdeck	PD §1-2	Sonnet 4.6	6,077
business-analyst-fullstory	FS §2	Sonnet 4.6	6,239
comp-eval-mkt-pos-pitchdeck	PD §3	Sonnet 4.6	4,931
comp-eval-moats-pitchdeck	PD §4	Sonnet 4.6	3,691
competitor-evaluator-fullstory	FS §3	Sonnet 4.6	6,023
financial-analyst-pitchdeck	PD §5,7,8	Sonnet 4.6	5,625
financial-analyst-fullstory	FS §6 Judge	Sonnet 4.6	4,197
management-evaluator-pitchdeck	PD §6	Sonnet 4.6	4,851
management-evaluator-fullstory	FS §4	Sonnet 4.6	6,004
risk-analyst-pitchdeck	PD §9	Sonnet 4.6	4,842
risk-analyst-fullstory	FS §1+Bear	Opus 4.6	7,443
valuation-specialist-pitchdeck	PD §10	Sonnet 4.6	5,090
valuation-specialist-fullstory	FS §5	Sonnet 4.6	5,169
synthesis-writer-pitchdeck	PD §11	Sonnet 4.6	4,145
synthesis-writer-fullstory	FS §6 Bull/Reb	Sonnet 4.6	4,041



4. Evaluation Framework

For each platform, the report scores on eight dimensions:

Dimension	What it measures
Multi-agent fit	First-class abstraction for orchestrator + specialists + waves
Parallelism within waves	Native parallel dispatch vs hand-rolled Promise.all
Failure handling	Retries, replay, checkpoints, supervisor patterns, durable execution
Multi-LLM support	Can we mix Claude / GPT-5 / Gemini per agent?
Migration ease	1 = rewrite required → 5 = drop in verbatim
Production maturity	Who runs it in prod? When did it GA?
Cost at our volume	~50-100 Pitch Decks + 20 Full Stories per month
Production parity	Reversible if Managed Agents callable_agents opens?



5. Platform Deep Dives

Eleven platforms across four tiers. Tier 1 = Anthropic's own self-hosted stack. Tier 2 = open-source multi-agent frameworks. Tier 3 = managed cloud agent services. Tier 4 = lower-level durable workflow engines you would combine with one of the above.

Tier 1 — Anthropic Self-Hosted Stack

5.1 Claude Agent SDK [TOP PICK]

The Claude Agent SDK (formerly Claude Code SDK) is Anthropic's open-source library that runs the same agent loop as Claude Code in your own infrastructure. Subagents are defined as Markdown files with YAML frontmatter — name, description, tools, model, permissions — almost identical to your current managed-agent.yaml + prompt.md split.

Property	Detail
Multi-agent	Subagents isolate context, run in parallel, parent dispatches via Task tool
Parallelism	Subagents fire in parallel by design; you control fan-out from dispatch code
Failure handling	You own retries (no managed supervisor); pair with Cloudflare Workflow for durability
Multi-LLM	Native Claude only; LiteLLM proxy enables OpenAI / Gemini / Bedrock
Migration	LOWEST. Drop prompts into .claude/agents/ with merged YAML frontmatter
Cost	SDK free; pay tokens (same as today) + ~\$10-25/mo runner hosting
Production parity	HIGHEST. Same model access, same prompts, easy revert
Gotchas	Re-implement session/event lifecycle; rewire web_search

Tier 2 — Multi-Agent Frameworks

5.2 LangGraph

The most production-tested multi-agent framework in market. Used by Klarna, Replit, Uber, LinkedIn, Elastic, JP Morgan, BlackRock. Hit v1.0 in Oct 2025. Agents are nodes in a StateGraph; edges are control flow. Execution is superstep-based (Pregel/BEAM model) — your wave structure maps 1:1 to graph layers.

- Parallelism: static fanout (multiple edges) or dynamic Send API (runtime map-reduce)
- Failure handling: per-node RetryPolicy, checkpointer (Postgres/SQLite/Dynamo), interrupt() for HITL
- Multi-LLM: trivial — ChatAnthropic, ChatOpenAI, ChatGoogleGenerativeAI interchangeable per node
- Cost: LangSmith Developer free (5K traces/mo); LangGraph self-hosted free for first 100K node executions
- Caveat: checkpoints are persistence, not durable execution — you need PM2/k8s/systemd to respawn the host

5.3 Inngest AgentKit



Best-fit AI-native abstraction for your specific pattern. Agent = single LLM + tools; Network = multiple agents + a Router that decides which agent runs next; backed by `step.run()` and `step.ai.infer()` which are auto-checkpointed and retried.

- Parallelism: native via `step.parallel()`; wave dispatch via Router function (~50 LOC)
- Failure handling: per-step retry, replay from last step on crash, throttling, concurrency limits
- Multi-LLM: native — OpenAI, Anthropic, Gemini, any OpenAI-compatible
- Cost: free tier 50K executions/mo (you live in free tier indefinitely)
- Production parity: lower than Claude Agent SDK; different abstraction

5.4 OpenAI Agents SDK

Production-ready open-source framework. Two coordination patterns: handoffs (control transfers) and agents-as-tools (orchestrator stays in control). Agents-as-tools fits your pipeline. LiteLLM adapter for non-OpenAI models is beta — for a Claude-heavy stack, Claude Agent SDK is a more direct fit.

5.5 Frameworks to skip

- CrewAI — hierarchical mode produces circular delegation, off-topic tangents, infinite consensus loops. Exact failure mode that would break your 11-agent dependency graph.
- Microsoft AutoGen v0.4 — conversation/group-chat mental model; does not match your fixed dependency-wave pipeline.
- Pydantic AI — promising AgentSpec for YAML loading; smaller ecosystem; revisit in 6 months.
- LlamaIndex AgentWorkflow — RAG-first lineage; multi-agent surface area is newer.

Tier 3 — Cloud Platform Agent Services

5.6 Cloudflare Project Think

The eventual right answer for Thes1s. April 2026 announcement adds sub-agents on Durable Objects, durable execution with fibers, persistent sessions, sandboxed code execution. AI Gateway brokers 70+ models with one API and no markup. Sub-agents are PREVIEW as of April 2026 — wait until GA before betting production on it.

5.7 AWS Bedrock Multi-Agent Collaboration

GA since early 2025, now part of Bedrock AgentCore. Hierarchical supervisor + collaborator agents. Sessions up to 8 hours. Claude Sonnet 4.6 GA in Bedrock Feb 2026 at same per-token price as Anthropic direct. Modest premium of \$0.10-0.50/run for AgentCore Runtime. Most mature multi-agent platform in market.

5.8 Vertex AI Agent Engine

ADK (Agent Development Kit) — Python+Java. Hierarchical agent trees, explicit `ParallelAgent` primitive. Free tier covers all your dev runs. Higher migration effort than Bedrock or LangGraph (~100 LOC per agent in Python).

5.9 Microsoft Foundry Agent Service



Best semantic match for waves (declarative State-based workflows; multi-agent workflows GA early 2026). Caveat: Enterprise/MCA-E subscriptions only — likely a blocker for an indie account.

Tier 4 — Durable Workflow Engines

Lower-level infrastructure. You would combine these with one of the agent frameworks above (or build a thin agent abstraction on top).

- Cloudflare Workflows — native to Workers; instances run for months; included in your \$5/mo plan. No agent abstraction. Best use: wrap a Claude Agent SDK runner in a Workflow step for durable execution.
- Temporal — gold standard. Most powerful, slowest migration, most ops burden. Cloud: \$50/M actions + \$100/mo base. Right answer at hedge-fund scale.
- Restate — newer Temporal-alike with simpler ops. Viable but no compelling edge over Inngest for your case.
- DBOS — durable execution as Postgres pattern. You don't have Postgres — pass.



6. Comparison Matrix

Scored 1-5 (5 = best for your specific Thes1s case). Higher total = better fit.

Platform	Multi	Parl	Fail	LLM	Mig	Mat	Cost	Par	Total
Inngest AgentKit	5	5	5	5	5	4	5	3	37
Claude Agent SDK	5	4	3	4	5	5	5	5	36
LangGraph	5	5	4	5	4	5	5	3	36
Microsoft Foundry*	5	5	4	4	4	4	5	3	34
AWS Bedrock MAC	5	5	5	3	3	5	4	3	33
CF Project Think	5	4	5	5	3	2	5	3	32
Vertex Agent Engine	4	5	4	5	3	4	4	3	32
OpenAI Agents SDK	4	4	3	4	3	5	5	2	30
Cloudflare Workflows	1	4	5	5	2	4	5	4	30
Temporal	1	5	5	5	1	5	3	4	29
CrewAI	3	3	2	5	2	4	4	2	25


* Foundry score gated by Enterprise subscription requirement.

Top three: Inngest AgentKit (37), Claude Agent SDK (36), LangGraph (36). All within striking distance — pick by deployment philosophy:

- Maximum production parity with Managed Agents → Claude Agent SDK
- Most AI-native abstraction with built-in durable execution → Inngest AgentKit
- Most mature framework with best observability story → LangGraph



Total Fit Score (out of 40)





7. Migration Effort by Platform

For each platform, what specifically you need to do to port the Pitch Deck pipeline. Time estimates assume one focused engineer.

Claude Agent SDK — ~5 days

Step	Days
Move prompt.md into .claude/agents/ with merged YAML frontmatter	1
Rewrite coordinator dispatch using Task tool	1
Re-wire web_search (Anthropic search API or Tavily)	0.5
Wrap pipeline in long-running Node service or Cloudflare Workflow	1
Port assembleDataPacket.js integration (mostly unchanged)	0.5
End-to-end test on LULU + 2 known-verdict tickers	1

Inngest AgentKit — ~5-6 days

Step	Days
Convert each agent to Agent({system, model, tools}) constructor	2
Write Router function for wave dispatch (~50 LOC)	1
Wire step.ai.infer() for each agent invocation	1
Sign up for Inngest, deploy initial run	0.5
End-to-end test	1

LangGraph — ~7 days

Step	Days
Define typed State TypedDict with reducers for parallel writes	1
Convert each agent to a node function	2
Wire StateGraph with explicit edges per wave	1
Set up Postgres checkpointer (or SQLite for dev)	0.5
Provision LangSmith for tracing	0.5
Stand up long-running runner (Fly.io / Render)	1
End-to-end test	1



AWS Bedrock MAC — ~7-10 days + AWS approval delay

CF Project Think — blocked on sub-agents GA (preview Apr 2026)




8. Cost Analysis

Assumes 100 Pitch Decks + 20 Full Stories per month. Token costs are LLM-side and identical across platforms (same Anthropic API prices). Platform fees are the differentiator.

Per-Run Cost (Pitch Deck)

Today, with Sonnet-only on Managed Agents, a Pitch Deck costs roughly \$5-6/run in tokens. The optimal mixed-model stack drops this to ~\$3.30/run.

Pitch Deck Cost Per Run (\$USD)



Monthly Platform Fees at Your Volume

Platform	Monthly Fee	Notes
Anthropic Managed Agents	\$0	Pay tokens only (today)
Claude Agent SDK	\$10-25	Hosting (Fly.io / Render) for runner
Inngest AgentKit	\$0	Free tier (50K executions/mo) covers all runs
LangGraph self-hosted	\$10-25	Hosting; LangSmith Dev free
LangGraph Cloud	\$10-20	Per run-minute pricing
AWS Bedrock MAC	\$15-50	AgentCore Runtime \$0.10-0.50/run
CF Project Think	\$0	Existing \$5/mo Workers plan
Vertex Agent Engine	\$0-100	Free tier covers dev; ~\$50-100 prod
Temporal Cloud	\$100+	Essentials base + per-action

Verdict: at your volume, platform fees are <\$50/mo for any reasonable choice. LLM token cost dominates and is roughly equal across platforms. Don't pick on platform fee — pick on capability fit.



9. Model Optimization Per Agent

This section applies regardless of which orchestration platform you pick. 19 of 20 agents on Sonnet is a missed opportunity. Mixing models per agent saves ~40% per Pitch Deck and lifts quality on numerical and adversarial sections.

Why Mix Models?

1. Specialization is real and benchmarked. GPT-5 leads FinanceReasoning at 88.23% (Opus 4.6 at 87.82%). Opus 4.6 leads multi-needle retrieval at 1M tokens (76% MRCR v2). Gemini 3.1 Pro tops Arena Creative Writing and abstract reasoning. No model wins everything.
2. Cost spread is 50x. Haiku 4.5 input: \$1/M tokens. Opus 4.7 output: \$25/M. Using Opus everywhere pays 25x for tasks Haiku does equally well.
3. Failure modes differ. GPT-5.4 reportedly refuses to fabricate missing financial data. Gemini 3.1 Pro is not recommended for incomplete-data financial tasks. Claude is most transparent about limitations.

Pricing Reference (April 2026)

Model	Input (\$/M)	Output (\$/M)	Context
Claude Opus 4.7	\$5	\$25	1M
Claude Sonnet 4.6	\$3	\$15	1M
Claude Haiku 4.5	\$1	\$5	200K
GPT-5 (standard)	\$1.25	\$10	400K
GPT-5.4 short-ctx	\$2.50	\$15	—
GPT-5.4 long-ctx	\$5.00	\$22.50	—
GPT-5 mini	\$0.25	\$2	—
Gemini 2.5 Pro ≤200K	\$1.25	\$10	1M
Gemini 2.5 Pro >200K	\$2.50	\$15	—

Recommended Per-Agent Assignment

Agent	Current	Proposed	Why	Per-Run \$
coordinator-pitchdeck	Sonnet 4.6	Haiku 4.5	JSON dispatch + routing, no reasoning	\$0.10
annual-reader	Sonnet 4.6	Gemini 2.5 Pro	MRCR retrieval parity at half price >200K	\$0.45
quarterly-reader	Sonnet 4.6	Gemini 2.5 Pro	Same; 10-Q + transcripts often >200K	\$0.30
business-analyst-pd	Sonnet 4.6	Sonnet 4.6 (keep)	Best balance long-form writing + cost	\$0.20
comp-mkt-position	Sonnet 4.6	GPT-5 (standard)	Numerical comparison; 88% FinanceReasoning	\$0.15



Agent	Current	Proposed	Why	Per-Run \$
comp-moats	Sonnet 4.6	Sonnet 4.6 (keep)	Web-search workflow; qualitative reasoning	\$0.20
financial-analyst-pd	Sonnet 4.6	GPT-5.4 short	Top FinanceReasoning; won't fabricate	\$0.40
management-eval-pd	Sonnet 4.6	GPT-5 (standard)	Mixed numerical (comp tables) + qualitative	\$0.15
risk-analyst-pd	Sonnet 4.6	Opus 4.7	Best Anthropic reasoning + extended thinking	\$0.50
valuation-specialist-pd	Sonnet 4.6	GPT-5.4 short	Highest accuracy on financial math	\$0.40
synthesis-writer-pd	Sonnet 4.6	Opus 4.7	Best long-form coherence with citation fidelity	\$0.45
risk-analyst-fs (Bear)	Opus 4.6	Opus 4.7	Adversarial reasoning + writing	\$0.40
synthesis-writer-fs	Sonnet 4.6	Opus 4.7 (Judge)	Final verdict needs highest reasoning	\$0.50

Sonnet-only baseline: ~\$5.50 per Pitch Deck. Mixed stack: ~\$3.30. Savings: ~40%, plus measurable quality lift on Valuation, Financial, and Risk.

The Production Parity Problem

feedback_production_parity says: don't make agent changes that won't translate to production Managed Agents. Today, Managed Agents only routes to Claude. So:

- Safe today (Anthropic-only swaps): Coordinator → Haiku 4.5, Risk Analyst PD → Opus 4.7, Synthesis Writer PD → Opus 4.7. Pure config changes that work in Managed Agents and any alternative platform.
- Cross-provider (Gemini, GPT) swaps: only viable if you switch off Managed Agents. This is the strongest pragmatic argument for the migration.

Phased Rollout

1. Phase A (no platform migration): Move Coordinator to Haiku 4.5; Risk PD to Opus 4.7; Synthesis PD to Opus 4.7. Estimated ~25% per-deck savings. 1 day.
2. Phase B (after platform migration): Long-context readers → Gemini 2.5 Pro. Financial Analyst + Valuation Specialist → GPT-5.4. Another ~15% savings. ~1 week of A/B testing across known-verdict tickers.
3. Phase C (calibration): Cross-model eval on ~10 known-verdict companies. Use Observatory wiki to track verdict-accuracy by model assignment.



10. Production Patterns We Are Missing

Modern multi-agent stacks come with capabilities that Managed Agents either doesn't expose or that you haven't wired up yet. Independent of platform choice. These raise the ceiling on what Thes1s can become.

Production Capabilities Stack (top = highest leverage)

Observability + Replay

LangSmith / Langfuse / AgentOps — trace, replay, time travel debug

Structured Eval Harness

LangSmith Datasets / OpenAI Evals / Promptfoo — known verdict pass/fail

Prompt Versioning

PromptHub / PromptLayer — measured impact per prompt change

Circuit Breakers + Cost Guards

Helicone / Portkey / AI Gateway — token & cost caps

Typed Output Validation

Pydantic AI / Zod / OpenAI strict mode — boundary enforcement

Cross-Run Memory Layer

Mem0 / Letta / AgentCore Memory — agent learns across tickers

A2A Agent Protocol

Standardized agent-to-agent for stickeR1 + 3rd party plug-ins

10.1 Observability with replay

LangSmith, Langfuse, AgentOps, Pydantic Logfire — all let you trace every LLM call across a multi-agent run, see token counts and latency per agent, and replay a failed run from the exact superstep that broke. Time-travel debugging lets you rewind to any checkpoint, modify state, and re-run downstream nodes. Single biggest dev-velocity multiplier for multi-agent work. Your Observatory wiki is a half-step in this direction.

10.2 Structured eval harness



Your Observatory has known-verdicts.json and verdict-check scripts. Modern eval harnesses let you define a dataset of known-verdict tickers as a fixture, run a candidate prompt/model assignment across the dataset in one command, get pass/fail + per-criterion scoring, and compare two configurations side-by-side with statistical significance. This is the engine for the RL-style optimization that the agent optimization phase needs.

10.3 Prompt versioning with measured impact

Right now your prompts live in git. With versioned prompts + datasets, you can answer: did my Annual Reader prompt change improve verdict accuracy on 10 calibration tickers? Today you have to manually run pipelines and eyeball verdict outcomes.

10.4 Circuit breakers and cost guards

Helicone / Portkey / AI Gateway — sit between your code and LLM APIs and enforce per-run token budget caps, per-month cost caps, automatic fallback to cheaper model on rate-limit, request deduplication and caching. For a hedge-fund-grade production system, these are not optional.

10.5 Typed structured output validation at boundaries

You already use Zod schemas for sections — that's the right pattern. Production-grade extension: validate at every agent boundary, retry with "Schema validation failed: <error>" appended to prompt when output is malformed. Pydantic AI typed agents and OpenAI strict mode make this enforceable at the API level.

10.6 Memory layer (cross-run learnings)

Mem0, Letta (formerly MemGPT), and the memory primitives in Claude Agent SDK + AgentCore Memory let agents accumulate cross-run knowledge. For Thes1s: an agent that processed AAPL last quarter remembers what surprised it, what it got wrong on the verdict check, and applies those learnings to MSFT. Real long-term competitive advantage.

10.7 A2A Agent Protocol

Google's A2A protocol (with broad industry adoption including Anthropic) lets agents from different vendors talk to each other in a standardized way. Important when Thes1s plugs into stickeR1's portfolio context, an external risk model, or a third-party valuation engine. Worth designing toward.



11. Recommended Path Forward

Step 1 — Continue waiting on Anthropic, but set a deadline

Email Anthropic again this week. Set an internal deadline of April 27. If callable_agents access hasn't landed by then, execute Step 2.

Step 2 — Execute Phase A model optimization NOW

Pure config change. Works on Managed Agents (One Pager) and any alternative.

- coordinator-pitchdeck → Haiku 4.5
- coordinator-fullstory → Haiku 4.5
- risk-analyst-pitchdeck → Opus 4.7
- synthesis-writer-pitchdeck → Opus 4.7 (final synthesis only)

Measure on the One Pager pipeline first. Run on 5 tickers. Compare verdict + cost vs current Sonnet baseline. Decide.

Step 3 — Migrate to Claude Agent SDK if Anthropic deadline misses

Allocate 5 days the week of April 28:

1. Stand up long-running Node service on Fly.io or Render
2. Move prompts to .claude/agents/ format
3. Rewrite Pitch Deck coordinator using Task tool dispatch
4. Wire web_search via Anthropic search API or Tavily
5. End-to-end test on LULU + 2 calibration tickers
6. Wrap in Cloudflare Workflow for durable execution

This preserves production parity. When Managed Agents callable_agents opens, you swap the dispatcher back without touching prompts.

Step 4 — Add observability (week 2)

Pick one: Langfuse (open source self-host), LangSmith (free tier 5K traces), or AgentOps. Wire as a passthrough on every LLM call. You will recoup the setup time on the first stuck pipeline run you can replay.

Step 5 — Run agent optimization sprint (weeks 3-4)

Now you have a platform that supports cross-provider routing. Execute Phase B: move readers to Gemini 2.5 Pro, Financial Analyst + Valuation Specialist to GPT-5.4. Run on all 10 known-verdict tickers. Use Observatory + tracing to track verdict accuracy per model assignment.



Step 6 — Long-term: Cloudflare Project Think (H2 2026)

Once Project Think exits preview, revisit. Your stack already lives on Cloudflare. This is the eventual home — single vendor, Durable Objects for state, AI Gateway for multi-LLM, all on infrastructure you understand.

Decision Tree

Decision Tree

```
graph TD
    Q1[Is Managed Agents callable_agents access live by Apr 27?]
    Q1 -- YES --> A1[stay on Managed Agents.  
Execute Phase A (Anthropic-only model swaps).  
Production parity preserved.]
    Q1 -- NO --> Q2[Add a 5-day engineering project this sprint?]
    Q2 -- YES --> A2[Claude Agent SDK [TOP PICK]  
closest to current architecture]
    Q2 -- NO --> A3[wait. Cost = One Pager-only revenue gap.]
    A3 --> Q3[Once migrated, want cross-provider model mixing?]
    Q3 -- YES --> A4[execute Phase B  
(Gemini for readers, GPT-5.4 for valuation)]
    Q3 -- NO --> A5[keep all-Anthropic, only Phase A]
    A5 --> Q4[Long-term Cloudflare-native rebuild?]
    Q4 -- YES --> A6[CF Project Think when sub-agents GA (H2 2026)]
    Q4 -- NO --> A7[stay on Claude Agent SDK indefinitely]
```



12. Appendix — Agent Inventory

Pitch Deck Agents (10 specialists + coordinator)

Agent	Wave	Tools
coordinator-pitchdeck	dispatcher	callable_agents, agent_toolset
annual-reader	Wave 0 (PSR)	filesystem
quarterly-reader	Wave 0 (PSR)	filesystem
business-analyst-pitchdeck	Wave 1	web_search
comp-eval-mkt-position-pitchdeck	Wave 1	web_search
comp-eval-moats-pitchdeck	Wave 2	web_search
financial-analyst-pitchdeck	Wave 2	web_search
management-evaluator-pitchdeck	Wave 2	web_search
risk-analyst-pitchdeck	Wave 3	web_search
valuation-specialist-pitchdeck	Wave 3	web_search
synthesis-writer-pitchdeck	Wave 4	synthesis only

Full Story Agents (7 specialists + coordinator)

Agent	Phase	Tools
coordinator-fullstory	dispatcher	callable_agents, agent_toolset
annual-reader	shared PSR	filesystem
quarterly-reader	shared PSR	filesystem
risk-analyst-fullstory	Phase 1 + Bear	web_search
business-analyst-fullstory	Phase 1	web_search
competitor-evaluator-fullstory	Phase 1	web_search
management-evaluator-fullstory	Phase 1	web_search
valuation-specialist-fullstory	Phase 1	web_search
financial-analyst-fullstory	Phase 2 (Judge)	synthesis only
synthesis-writer-fullstory	Phase 2 (Bull/Reb)	synthesis only

Source References

- Multi-agent: LangGraph docs, Inngest AgentKit, Claude Agent SDK Subagents, OpenAI Agents SDK



- Cloud platforms: AWS Bedrock MAC GA, CF Project Think, Vertex Agent Engine, MS Foundry Agent Service
- Durable execution: Temporal AI workflows, CF Workflows GA, Diagrid checkpoints critique
- Pricing: LangSmith, Inngest, Bedrock AgentCore, Anthropic, OpenAI, Vertex AI
- Model benchmarks: LM Council Apr 2026, Patronus FinanceBench, EQ-Bench Longform, Vellum Leaderboard, Chroma Context Rot